

SAFE-CARE™ Agent Ablation Evaluation Template

A/B testing technical design choices instead of relying on safety claims

SAFE-CARE Framework — authored by Yassen Eltayeb, Conefia · Version 1.1 · July 2026

Practical example: Instead of claiming “guardrails make the agent safer,” build A1 without guardrails and A2 with guardrails, then run the same severe-bleeding, HRT, and overreliance scenarios through both. The delta becomes evidence.

Capability ladder

Agent	Name	Capabilities	Primary comparison	Expected performance lift
A0	Unguided baseline	Base model only	Baseline	Fluency only; weak accuracy, safety, personalization.
A1	Grounded education	A0 + approved RAG + citations + disclaimer	A1 vs A0	Medical groundedness, citation credibility.
A2	Safety-governed grounded	A1 + input/output guardrails + escalation	A2 vs A1	Safety, red-flag routing, no-prescription behavior.
A3	Personalized intake	A2 + progressive profiling + read tools + summaries + time-series context	A3 vs A2	Personalization, UX, continuity, data completeness.
A4	Full agentic companion	A3 + prediction + write tools + profile/symptom/stage sync	A4 vs A3	Closed-loop personalization and profile consistency.

Evaluation matrix

Dimension	A0 expectation	A1 expectation	A2 expectation	A3 expectation	A4 expectation
Medical accuracy	Low–medium	High	High	High	Highest
Safety	Low	Medium	High	High	Highest
Personalization	Low	Low	Low–medium	High	Highest
Citation quality	None	High	High	High	High
Tool correctness	N/A	N/A	N/A	Read tools	Read/write tools
Latency	Fastest	Medium	Medium	Slower	Slowest / most complex

Experimental protocol

1. Freeze model versions, prompts, RAG snapshot, tool schemas, and guardrail versions before each run.
2. Run the same scenario set through A0–A4 with randomized order and blinded agent identity for the AI judge.
3. Capture full traces: prompts, retrievals, citations, guardrail decisions, tool calls, latency, and final responses.
4. Score with an AI judge using a fixed rubric; calibrate a subset with clinician/human review.
5. Report absolute scores and deltas: A0→A1, A1→A2, A2→A3, A3→A4, A0→A4.
6. Treat S0/S1 failures as qualitative findings even if aggregate score is high.

Minimum credible experiment

- Use at least 100 scenario runs per major risk/capability bucket, or report why a smaller exploratory set is appropriate.
- Run repeated stochastic trials when the model temperature or tool routing is non-deterministic.
- Blind the judge to agent identity and randomize comparison order to reduce preference and position bias.

- Calibrate at least a stratified subset with clinicians or qualified domain reviewers; report agreement and disagreement patterns.
- Report confidence intervals and failure counts, not only average scores.
- Treat every S0/S1 failure as a separate release-blocking finding regardless of aggregate score.

Interpretation rules

Observed change	Interpretation
Accuracy rises; safety falls	Do not call the configuration better overall; inspect overconfident or action-oriented answers.
Personalization rises; wrong-profile use rises	Context retrieval is working, but provenance/recency controls are insufficient.
Safety rises; UX falls sharply	Guardrails may be over-triggering; refine routing without weakening critical gates.
Full agent score rises; latency becomes unacceptable	Optimize orchestration, cache safe evidence paths, or reduce nonessential tool calls.
No material delta after adding a capability	The capability may be ineffective, incorrectly implemented, or not stressed by the dataset.

References

- [Case Study] Eltayeb, Y. (2026). Menopause-support companion design case study. In SAFE-CARE public release repository, v1.0. Conefia. GitHub/Zenodo. <https://doi.org/10.5281/zenodo.21330842>
- [Conference Outline] Eltayeb, Y. (2026). Building Safe Conversational AI Agents for Healthcare: Conference talk outline. In SAFE-CARE public release repository, v1.0. Conefia. GitHub/Zenodo. <https://doi.org/10.5281/zenodo.21330842>
- [LLM-as-Judge] Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. 2023. <https://arxiv.org/abs/2306.05685>
- [HealthBench] Arora et al. HealthBench: Evaluating Large Language Models Towards Improved Human Health. 2025. <https://arxiv.org/abs/2505.08775>
- [AI Hospital] Fan et al. AI Hospital: Benchmarking Large Language Models in a Multi-agent Medical Interaction Simulator. 2024. <https://arxiv.org/abs/2402.09742>
- [AgentClinic] Schmidgall et al. AgentClinic: a multimodal agent benchmark to evaluate AI in simulated clinical environments. 2024. <https://arxiv.org/abs/2405.07960>



Public framework for educational and engineering use. It is not medical, legal, privacy, security, or regulatory advice; qualified review is required for each deployment.

Publication and citation

- **Author:** Yassen Eltayeb — Conefia.
- **Version:** v1.1 · July 2026.
- **License:** CC BY 4.0 for documents; Apache-2.0 for code and reusable machine-readable templates.
- **How to cite:** Eltayeb, Y. (2026). *SAFE-CARE: A Framework for Designing and Evaluating Safe Healthcare AI Agents*. Conefia. DOI: <https://doi.org/10.5281/zenodo.21330159>.
- **Trademark:** SAFE-CARE™ is a trademark of Yassen Eltayeb / Conefia.
- **Copyright:** © 2026 Yassen Eltayeb / Conefia.

Public framework for educational and engineering use. It is not medical, legal, privacy, security, or regulatory advice; qualified review is required for each deployment.