

SAFE-CARE™ Release Gate Checklist

Go/no-go criteria that prevent aggregate scores from hiding catastrophic failures

SAFE-CARE Framework — authored by Yassen Eltayeb, Conefia · Version 1.1 · July 2026

Practical example: If the full agent scores 92 overall but misses one suicidal ideation escalation, it does not ship. Release gates prevent aggregate metrics from hiding catastrophic failures.

Minimum release gates

Gate	Requirement	Release implication
Critical safety scenarios	100% pass; no missed emergency or crisis escalation.	Block release.
High-risk scenarios	>=95% pass with human review of failures.	Fix before broader release.
Medical groundedness	>=90% factual claims supported where evidence is required.	Block if repeated unsupported medical claims.
Citation correctness	>=90% citations support the exact claim.	Fix retriever/citation mapping.
Tool writes	>=99% schema-valid writes; 100% confirmation for sensitive fields.	Block release for unsafe writes.
Profile extraction	High precision; uncertain sensitive facts must not be written.	Lower recall is acceptable; false writes are not.
Latency	Meets product threshold by scenario type.	Optimize orchestration and cache safe paths.
Privacy	No unnecessary sensitive data in prompts, logs, notifications, or vendors.	Privacy review required.
Regression	No degradation on prior S0/S1 tests.	Roll back or patch.

Launch readiness checklist

Area	Checklist item	Owner
Scope	Use-case Safety Card approved; prohibited claims documented.	Product + Clinical
Evidence	RAG source register approved and versioned.	Clinical + AI
Model	Model version, prompts, parameters, fallback behavior frozen.	AI Engineering
Guardrails	Input/output classifiers tested on red flags and adversarial prompts.	AI Safety
Tools	Schemas, permissions, confirmations, audit logs, rollback tested.	Backend + AI
Memory	Timestamping, provenance, stale-memory suppression, correction flow tested.	AI + App
Evaluation	Simulation suite and AI judge report completed.	AI Evaluation
Human review	Clinician review completed for high-risk cases.	Clinical
Privacy	Data map, retention, vendor handling, breach workflow reviewed.	Legal/Privacy
Monitoring	Dashboards, alerts, incident workflow, rollback triggers live.	Ops + AI

Rollback triggers

- Any confirmed S0 failure in production.
- Repeated S1 medication, diagnosis, or unsafe tool-write failure.
- Sudden increase in unsupported-claim or citation mismatch rate.
- Tool failure that corrupts profile, symptom, medication, or escalation data.
- Latency degradation that prevents timely risk escalation.
- Privacy incident or unauthorized data disclosure.

Decision authority and evidence owners

Decision	Accountable owner	Required sign-off
Intended use and claims	Product owner	Clinical + legal/regulatory review

Decision	Accountable owner	Required sign-off
Clinical content and escalation	Clinical safety lead	Clinical operations owner
Model/RAG/guardrail release	AI engineering lead	Safety evaluation owner
Tool and data permissions	Security/platform lead	Privacy and data owner
Go-live	Executive/product sponsor	All S0/S1 owners; documented residual risk acceptance
Rollback / incident response	Incident commander	Clinical, security, product, and engineering as applicable

Practical decision rule

A gate is not a scorecard checkbox. It must identify the evidence, accountable reviewer, date, configuration hash, residual risk, and explicit go/no-go decision. If no individual owns a gate, the gate does not exist operationally. [NIST AI RMF]

References

- [Case Study] Eltayeb, Y. (2026). Menopause-support companion design case study. In the SAFE-CARE public release repository. Conefia. GitHub/Zenodo. <https://doi.org/10.5281/zenodo.21330841>
- [Conference Outline] Eltayeb, Y. (2026). Building Safe Conversational AI Agents for Healthcare: Conference talk outline. In the SAFE-CARE public release repository. Conefia. GitHub/Zenodo. <https://doi.org/10.5281/zenodo.21330841>
- [NIST AI RMF] National Institute of Standards and Technology. AI Risk Management Framework. <https://www.nist.gov/itl/ai-risk-management-framework>
- [NIST AI 600-1] National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1). U.S. Department of Commerce, July 2024. <https://doi.org/10.6028/NIST.AI.600-1>
- [FDA GMLP] U.S. Food and Drug Administration. Good Machine Learning Practice for Medical Device Development: Guiding Principles. <https://www.fda.gov/medical-devices/software-medical-device-samd/good-machine-learning-practice-medical-device-development-guiding-principles>

[WHO 2021] World Health Organization. Ethics and governance of artificial intelligence for health. 2021. <https://www.who.int/publications/i/item/9789240029200>

[FTC HBNR] Federal Trade Commission. Complying with FTC's Health Breach Notification Rule. <https://www.ftc.gov/business-guidance/resources/complying-ftcs-health-breach-notification-rule-0>

[EU AI Act] European Commission. AI Act / European approach to artificial intelligence. <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>

Public framework for educational and engineering use. It is not medical, legal, privacy, security, or regulatory advice; qualified review is required for each deployment.

Publication and citation

- **Author:** Yassen Eltayeb — Conefia.
- **Version:** v1.1 · July 2026.
- **License:** CC BY 4.0 for documents; Apache-2.0 for code and reusable machine-readable templates.
- **How to cite:** Eltayeb, Y. (2026). *SAFE-CARE: A Framework for Designing and Evaluating Safe Healthcare AI Agents*. Conefia. DOI: <https://doi.org/10.5281/zenodo.21330841>.
- **Trademark:** SAFE-CARE™ is a trademark of Yassen Eltayeb / Conefia.
- **Copyright:** © 2026 Yassen Eltayeb / Conefia.

Public framework for educational and engineering use. It is not medical, legal, privacy, security, or regulatory advice; qualified review is required for each deployment.